

Reinforcement Learning with Deep Energy-Based Policies

Tuomas Haarnoja, Haoran Tang, Pieter Abbeel, Sergey Levine

ICML2017

May 14, 2019

Outline

Introduction

Preliminaries

- Maximum Entropy RL
- Energy-Based Models

Soft Q-Learning

- Integrating problem
- Sample problem

Algorithm

Experiments

- Multi-Goal Environment
- Task-specific initialization

Conclusion

Introduction

Deep RL training is sensitive to ...

- ▶ randomness in the environment
- ▶ initialization of the policy
- ▶ the algorithm implementation

Example




Figure 1: A robot navigating a maze

Introduction

RL employs a (stochastic) policy (π) to select actions, and finds the best policy that maximizes the cumulative reward it collects throughout an episode of length T .

$$\pi^* = \operatorname{argmax}_{\pi} \mathbb{E}_{\pi} \left[\sum_{t=0}^T r_t \right]$$




Figure 2: A multimodal Q-function.

Introduction

Energy Based Model and Boltzmann Distribution and Softmax

- ▶ Energy Based Model: $P(x) = \frac{e^{f(x)}}{\sum_{x \in X} e^{f(x)}}$
- ▶ Boltzmann Distribution : $P(x) = \frac{e^{-E(x)}}{\sum_{x \in X} e^{-E(x)}}$
- ▶ Softmax function : $\sigma(\mathbf{z})_j = \frac{e^{z_j}}{\sum_{k=1}^K e^{z_k}}$

In practice, we prefer maximum-entropy models as they assume the least about the unknowns while matching the observed information.

Preliminaries

Maximum Entropy RL

the state and state-action marginals of the trajectory distribution induced by a policy $\pi(a_t|s_t)$ are:

$$\rho_{\pi}(\mathbf{s}_t) \text{ and } \rho_{\pi}(\mathbf{s}_t, \mathbf{a}_t)$$

the standard RL objective is:

$$\pi_{\text{std}}^* = \arg \max_{\pi} \sum \mathbb{E}_{(\mathbf{s}_t, \mathbf{a}_t) \sim \rho_{\pi}} [r(\mathbf{s}_t, \mathbf{a}_t)]$$

the maximum entropy RL objective is:

$$\pi_{\text{MaxEnt}}^* = \arg \max_{\pi} \sum_t \mathbb{E}_{(\mathbf{s}_t, \mathbf{a}_t) \sim \rho_{\pi}} [r(\mathbf{s}_t, \mathbf{a}_t) + \alpha \mathcal{H}(\pi(\cdot|\mathbf{s}_t))]$$

Preliminaries

Energy-Based Models

soft Q-function is:

$$Q_{\text{soft}}^* (\mathbf{s}_t, \mathbf{a}_t) = r_t + \mathbb{E}_{(\mathbf{s}_{t+1}, \dots) \sim \rho_\pi} \left[\sum_{l=1}^{\infty} \gamma^l (r_{t+l} + \alpha \mathcal{H}(\pi_{\text{MaxEnt}}^* (\cdot | \mathbf{s}_{t+l}))) \right]$$

and it satisfies the soft Bellman equation:

$$Q_{\text{soft}}^* (\mathbf{s}_t, \mathbf{a}_t) = r_t + \gamma \mathbb{E}_{\mathbf{s}_{t+1} \sim p_s} [V_{\text{soft}}^* (\mathbf{s}_{t+1})],$$

where the soft value function is:

$$V_{\text{soft}}^* (\mathbf{s}_t) = \alpha \log \int_{\mathcal{A}} \exp \left(\frac{1}{\alpha} Q_{\text{soft}}^* (\mathbf{s}_t, \mathbf{a}') \right) d\mathbf{a}'$$

the optimal policy is:

$$\pi_{\text{MaxEnt}}^* (\mathbf{a}_t | \mathbf{s}_t) = \exp \left(\frac{1}{\alpha} (Q_{\text{soft}}^* (\mathbf{s}_t, \mathbf{a}_t) - V_{\text{soft}}^* (\mathbf{s}_t)) \right)$$

Soft Q-Learning

q-learning: $Q(s_t, a_t) \leftarrow r_t + \gamma \max_a Q(s_{t+1}, a)$

soft q-learning: $Q_{\text{soft}}(\mathbf{s}_t, \mathbf{a}_t) \leftarrow r_t + \gamma \mathbb{E}_{\mathbf{s}_{t+1} \sim p_s} [V_{\text{soft}}(\mathbf{s}_{t+1})], \forall \mathbf{s}_t, \mathbf{a}_t$
 $V_{\text{soft}}(\mathbf{s}_t) \leftarrow \alpha \log \int_{\mathcal{A}} \exp\left(\frac{1}{\alpha} Q_{\text{soft}}(\mathbf{s}_t, \mathbf{a}')\right) d\mathbf{a}', \forall \mathbf{s}_t$

In fact, the q-learning uses hard-max of Q, the soft q-learning uses soft-max of Q.

In continuous domains, there are two major challenges of soft-q-learning.

- ▶ First, exact dynamic programming is infeasible
- ▶ Second, the optimal policy is defined by an intractable energy-based distribution, which is difficult to sample from.

Soft Q-Learning

Integrating problem

we use samples from rollouts of the current policy, and use sample sum instead of integrating.

- ▶ Model the soft Q-function with a function approximator, and denote it as $Q_{\text{soft}}^{\theta}(\mathbf{s}_t, \mathbf{a}_t)$
- ▶ Express the soft value function in terms of an expectation via importance sampling:

$$V_{\text{soft}}^{\theta}(\mathbf{s}_t) = \alpha \log \mathbb{E}_{q_{\mathbf{a}'}} \left[\frac{\exp \left(\frac{1}{\alpha} Q_{\text{soft}}^{\theta}(\mathbf{s}_t, \mathbf{a}') \right)}{q_{\mathbf{a}'}(\mathbf{a}')} \right]$$

- ▶ Minimize the objective:

$$J_Q(\theta) = \mathbb{E}_{\mathbf{s}_t \sim q_{\mathbf{s}_t}, \mathbf{a}_t \sim q_{\mathbf{a}_t}} \left[\frac{1}{2} \left(\hat{Q}_{\text{soft}}^{\bar{\theta}}(\mathbf{s}_t, \mathbf{a}_t) - Q_{\text{soft}}^{\theta}(\mathbf{s}_t, \mathbf{a}_t) \right)^2 \right]$$

Soft Q-Learning

Sample problem

We use the amortized Stein variational gradient descent (SVGD) to train an inference network to generate approximate samples

- ▶ Construct a state-conditioned stochastic neural network
 $\mathbf{a}_t = f^\phi(\xi; \mathbf{s}_t)$
- ▶ Minimize the distance between the induced distribution and the energy-based distribution:

$$J_\pi(\phi; \mathbf{s}_t) = D_{\text{KL}}(\pi^\phi(\cdot | \mathbf{s}_t) \parallel \exp\left(\frac{1}{\alpha} (Q_{\text{soft}}^\theta(\mathbf{s}_t, \cdot) - V_{\text{soft}}^\theta)\right))$$

Soft Q-Learning

We can use the chain rule and backpropagate the Stein variational gradient into the policy network according to

$$\frac{\partial J_{\pi}(\phi; \mathbf{s}_t)}{\partial \phi} \propto \mathbb{E}_{\xi} \left[\Delta f^{\phi}(\xi; \mathbf{s}_t) \frac{\partial f^{\phi}(\xi; \mathbf{s}_t)}{\partial \phi} \right]$$

where $\Delta f^{\phi}(\cdot; \mathbf{s}_t)$ is Stein variational gradient descent(SVGD):

$$\begin{aligned} \Delta f^{\phi}(\cdot; \mathbf{s}_t) = & \mathbb{E}_{\mathbf{a}_t \sim \pi^{\phi}} \left[\kappa(\mathbf{a}_t, f^{\phi}(\cdot; \mathbf{s}_t)) \nabla_{\mathbf{a}'} Q_{\text{soft}}^{\theta}(\mathbf{s}_t, \mathbf{a}') \Big|_{\mathbf{a}'=\mathbf{a}_t} \right. \\ & \left. + \alpha \nabla_{\mathbf{a}'} \kappa(\mathbf{a}', f^{\phi}(\cdot; \mathbf{s}_t)) \Big|_{\mathbf{a}'=\mathbf{a}_t} \right], \end{aligned}$$

and κ is a kernel function.

Besides, we can use any gradient-based optimization method to learn the optimal sampling network parameters.

Algorithm

Algorithm 1 Soft Q-learning

$\theta, \phi \sim$ some initialization distributions.

Assign target parameters: $\bar{\theta} \leftarrow \theta, \bar{\phi} \leftarrow \phi$.

$\mathcal{D} \leftarrow$ empty replay memory.

for each epoch **do**

for each t **do**

Collect experience

 Sample an action for \mathbf{s}_t using f^ϕ :

$\mathbf{a}_t \leftarrow f^\phi(\xi; \mathbf{s}_t)$ where $\xi \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$.

 Sample next state from the environment:

$\mathbf{s}_{t+1} \sim p_{\mathbf{s}}(\mathbf{s}_{t+1} | \mathbf{s}_t, \mathbf{a}_t)$.

 Save the new experience in the replay memory:

$\mathcal{D} \leftarrow \mathcal{D} \cup \{(\mathbf{s}_t, \mathbf{a}_t, r(\mathbf{s}_t, \mathbf{a}_t), \mathbf{s}_{t+1})\}$.

Sample a minibatch from the replay memory

$\{(\mathbf{s}_t^{(i)}, \mathbf{a}_t^{(i)}, r_t^{(i)}, \mathbf{s}_{t+1}^{(i)})\}_{i=0}^N \sim \mathcal{D}$.

Update the soft Q-function parameters

 Sample $\{\mathbf{a}^{(i,j)}\}_{j=0}^M \sim q_{\mathbf{a}'} \text{ for each } \mathbf{s}_{t+1}^{(i)}$.

 Compute empirical soft values $\hat{V}_{\text{soft}}^{\bar{\theta}}(\mathbf{s}_{t+1}^{(i)})$ in (10).

 Compute empirical gradient $\hat{\nabla}_{\theta} J_Q$ of (11).

 Update θ according to $\hat{\nabla}_{\theta} J_Q$ using ADAM.

Update policy

 Sample $\{\xi^{(i,j)}\}_{j=0}^M \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ for each $\mathbf{s}_t^{(i)}$.

 Compute actions $\mathbf{a}_t^{(i,j)} = f^\phi(\xi^{(i,j)}, \mathbf{s}_t^{(i)})$.

 Compute Δf^ϕ using empirical estimate of (13).

 Compute empirical estimate of (14): $\hat{\nabla}_{\phi} J_{\pi}$.

 Update ϕ according to $\hat{\nabla}_{\phi} J_{\pi}$ using ADAM.

end for

if epoch \bmod update_interval = 0 **then**

 Update target parameters: $\bar{\theta} \leftarrow \theta, \bar{\phi} \leftarrow \phi$.

end if

end for

Experiments

Multi-Goal Environment




Figure 3: Illustration of the 2D multi-goal environment

Experiments

Task-specific initialization

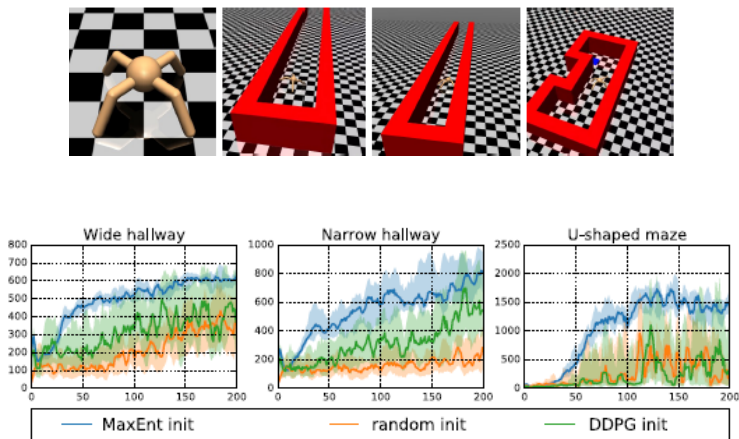


Figure 4: Performance in the downstream task with fine-tuning (MaxEnt) or training from scratch (DDPG).

Conclusion

- ▶ Energy-based Model and Softmax
- ▶ Use Importance Sampling to deal with Integrating problem
- ▶ Construct sample network

Thanks